Yuno code benchmark v4 — nine models, one baseline plus two late entries

AI Squad · Yuno
2026-09-28

> **Verdict.** Seven models regenerated on the same day, same protocol, same rebuilt suite, form the baseline; **Claude Opus 5.5 (2026-09-22) and Claude Sonnet 5.5 (2026-09-28) were added via Amazon Bedrock** (same suite and protocol, different day and provider). **Claude Sonnet 5.5 passes 48/48 with no refusals**, the first Claude model to clear the suite, at 8.5 s per problem. **Grok 4.5 and Gemini 3.8 Flash also pass 48/48**; Gemini is the cheapest run ($0.72) and twice as fast as Grok. **Opus 5.5 passes 46/48** (Opus 5: 35/48) — with 2 content-filter refusals and 100% correct on everything it attempts. **Claude Fable 5.1 passes 40/48**: it refuses 8 prompts (optimization and algorithms) and is also 100% correct on what it attempts. Kimi K3 (40) and GLM 5.2 (35) lose answers by exhausting the 16k output cap on reasoning; Opus 5 and Fable 5 (35 each) refuse 13 and are perfect on the rest. The Claude raw-API refusal rate remains the operational risk to watch, but it has moved from 30/48 (Fable 5, July) to 8/48 (Fable 5.1) to 2/48 (Opus 5.5) and 0/48 (Sonnet 5.5).

What was tested

48 in-house HumanEval-style Python problems in six categories (algorithms, data structures, optimization, bug fixing, spec compliance, edge cases; 8 each). Each model receives the problem prompt as the only user message, no system prompt (Claude: no temperature; OpenRouter models: temperature 0). Grading is by execution: the candidate module is imported in a sandboxed child process (memory and time limits) and run against the test cases; `pass` requires every case, `tests_ratio` is the fraction passed, `checker` is 1 when the module loads and defines the entry point. Refusals (`stop_reason: refusal`) and empty answers score 0 and are reported separately so they are never confused with wrong code.

The v4 suite. Both July runners were lost, so the suite was rebuilt from the problems and canonical solutions stored in Braintrust: 461 cases (6-15 per problem), expected values generated by executing the canonical solution, and a large deterministic input per optimization problem sized so a quadratic solution times out while the canonical runs in under a second. The suite was calibrated against the stored July 25 outputs of five models: 89.6% agreement on pass (215/240). Every disagreement was reviewed: 21 go in the model's favour (July tests stricter than the prompt; the outputs are behaviourally identical to the canonical under fuzzing) and 4 are real bugs July missed. Candidates run on Python 3.12 (the prompts never pin a version; on 3.9 two correct Gemini solutions failed on `int | None` and `bisect(key=)`). The suite, harness and every output now live in the `code-benchmarks` repository.

One baseline, plus two late entries. The seven-model baseline was generated on September 2 with a 16,384-token output cap. Claude Opus 5.5 was added separately on September 22 via Amazon Bedrock (inference profile `us.anthropic.claude-opus-5-5`, us-east-1) and Claude Sonnet 5.5 on September 28 (inference profile `global.anthropic.claude-sonnet-5-5`, us-east-1; Bedrock offers no

`us.` profile for it), both with the same 16,384-token cap and prompt protocol — not part of the single-day baseline, but scored on the identical v4 suite. Two kinds of empty answer are reported separately: the Anthropic classifier refusing a benign prompt (`stop_reason: refusal` on the direct API, `stopReason: content_filtered` on Bedrock — normalized to the same label) and a reasoning model exhausting the output cap without emitting code (`finish_reason: length`; retrying at 32k changed nothing, the provider caps Kimi and GLM at 16,384). Two Opus 5 requests hit 529 Overloaded and were retried successfully. Bedrock reports no per-call cost, so the Opus 5.5 and Sonnet 5.5 cost columns read n/a rather than a list-price estimate.

Results

MODEL	PASS (RAW)	WHEN ATTEMPTED	EMPTY	LATENCY	OUTPUT TOKENS	COST / RUN	SOURCE
Grok 4.5	48/48 (100%)	100%	0	45.8 s	2,586	$0.79	OpenRouter
Gemini 3.8 Flash	48/48 (100%)	100%	0	21.7 s	3,965	$0.72	OpenRouter
Claude Sonnet 5.5	48/48 (100%)	100%	0	8.5 s	1,117	n/a	Bedrock
Claude Opus 5.5	46/48 (96%)	100%	2 refusals	16.1 s	1,528	n/a	Bedrock
Claude Fable 5.1	40/48 (83%)	100%	8 refusals	17.0 s	1,263	≈$2.83	Anthropic API
Kimi K3	40/48 (83%)	95%	6 output cap	40.5 s	3,909	$4.09	OpenRouter
GLM 5.2	35/48 (73%)	97%	12 output cap	27.3 s	5,465	$1.99	OpenRouter
Claude Opus 5	35/48 (73%)	100%	13 refusals	23.6 s	1,978	≈$1.88	Anthropic API
Claude Fable 5	35/48 (73%)	100%	13 refusals	18.5 s	1,507	≈$2.95	Anthropic API

Empty = API refusal (Claude 5) or output cap exhausted by reasoning (Kimi, GLM). "When attempted" = passes ÷ answered prompts. Latency and tokens are per-problem means on answered prompts. Cost per 48-problem run: measured by OpenRouter; ≈ list-price estimate (tokens × rate) for the direct Anthropic API.

By category (passes out of 8)

CATEGORY	GROK 4.5	GEMINI 3.8 FLASH	SONNET 5.5	OPUS 5.5	FABLE 5.1	KIMI K3	GLM 5.2	OPUS 5	FABLE 5
Algorithms	8/8	8/8	8/8	7/8 (1×)	5/8 (3×)	5/8 (3×)	5/8 (3×)	4/8 (4×)	4/8 (4×)
Data structures	8/8	8/8	8/8	8/8	8/8	7/8 (1×)	5/8 (3×)	8/8	8/8
Optimization	8/8	8/8	8/8	7/8 (1×)	3/8 (5×)	8/8	5/8 (3×)	3/8 (5×)	1/8 (7×)
Bug fixing	8/8	8/8	8/8	8/8	8/8	6/8 (1×)	7/8 (1×)	4/8 (4×)	6/8 (2×)
Spec compliance	8/8	8/8	8/8	8/8	8/8	8/8	6/8 (1×)	8/8	8/8
Edge cases	8/8	8/8	8/8	8/8	8/8	6/8 (1×)	7/8 (1×)	8/8	8/8

July 25 outputs vs September 2 outputs, both scored on v4

MODEL	JUL 25 GENERATION	SEP 2 GENERATION	Δ
Grok 4.5	45/48	48/48	+3
Kimi K3	47/48	40/48	-7
GLM 5.2	44/48	35/48	-9
Claude Opus 5	35/48	35/48	+0
Claude Fable 5	16/48	35/48	+19

Same suite, same scoring; only the generation date differs. Fable 5 gains 19 problems purely from fewer refusals (30 → 13) on identical prompts; Kimi and GLM lose problems to longer reasoning that now exhausts the output cap; Grok fixes the two cases its July code got wrong.

Findings

- **Grok 4.5 and Gemini 3.8 Flash clear the whole suite.** Both 48/48 with no empty answers. Gemini does it at $0.72 and 21.7 s per problem (3965 output tokens); Grok takes 45.8 s per problem, twice its July pace, with longer answers (2.6k tokens). Grok also fixes the two cases its July outputs got wrong (FIFO scheduler re-submission, `str.format` where forbidden).
- **Sonnet 5.5 clears the suite, fastest of all nine models.** 48/48, no refusals, 8/8 in every category, at 8.5 s and 1,117 output tokens per problem — faster than Gemini 3.8 Flash and roughly half the time of Opus 5.5. It is the first Claude model with zero refusals on these bare code prompts.
- **Opus 5.5 is close behind.** 46/48 (2 content-filter refusals, one algorithms and one optimization problem) and 100% correct on what it attempts, at 1,528 output tokens and 16.1 s per problem — shorter and faster than Opus 5, and a jump from Opus 5's 35/48. Its refusal rate (2/48) is also the lowest of any Claude generation tested, direct API or Bedrock.
- **Fable 5.1 attempts far more and never misses.** 8 refusals (5 optimization, 3 algorithms; the other four categories are 8/8) and 40/40 on what it answers, with the shortest answers in the set (1263 tokens, 17.0 s). Fable 5 itself, regenerated today, refuses 13 instead of July's 30 on identical prompts: the classifier moved on Anthropic's side, and the model is 35/35 when it answers. Opus 5: 13 refusals, 35/35.
- **Reasoning models hit the output cap.** Kimi K3 leaves 6 problems and GLM 5.2 leaves 12 with no code after 16,384 output tokens of reasoning; retrying with a 32k budget changed nothing because the provider caps both there. On what they deliver they are strong (Kimi 40/42, GLM 35/36), but Kimi is now the most expensive run ($4.09) and, with Grok, the slowest.
- **Two empties, one score.** A refusal and an exhausted cap both score 0 like a wrong answer and are invisible unless `stop_reason` / `finish_reason` (or Bedrock's `content_filtered`) is inspected. Every Claude refusal in this suite came from bare code prompts over an API with no system prompt; the Maia agent evals run on Fable without a single refusal.

Recommendation

For autonomous code agents on the raw API, Gemini 3.8 Flash and Grok 4.5 are the frictionless choices; Gemini wins on cost and latency. Where a Claude model is the requirement, Sonnet 5.5 is now the pick for code on this suite — full marks, no refusals, the fastest run, and available on Bedrock without an Anthropic API key; Opus 5.5 is a close second. Note that for Maia's investigation work the picture differs: in the same-day agent evals Fable 5.1 beat Sonnet 5.5 (93% vs 82% adjusted), so the right model depends on the job. Fable 5.1 remains viable where a supervising loop can detect a refusal and re-route (it is 100% correct on what it attempts), and refusal rates should be re-measured on every release: Fable 5 went from 30 to 13 refusals in five weeks with no version change, and Opus 5.5 is down to 2. Give reasoning models (Kimi, GLM) a larger output budget where the provider allows it, or expect empty answers on the hardest prompts. Keep the suite in the repository and re-run `calibrate` before every comparison.

Appendix — non-refusal failures per model

MODEL	CATEGORY	PROBLEM	FIRST FAILING CASE
Kimi K3	bug_fixing	`recursive-settings-merger-drops-falsy-overrides`	import: NameError: name 'merge_settings' is not defined
Kimi K3	edge_cases	`even-chunking-of-arbitrary-iterables`	import: NameError: name 'chunk_evenly' is not defined
GLM 5.2	spec_compliance	`top-k-frequent-without-sorting-builtins`	no_sorting_builtins: expected {"value": [false, false, false, false, false]} got {"ok": true, "loaded": true, "value": [false, true, false, false, fal

Sources: Braintrust project "code-benchmarks", experiments code-bench-v4__* (seven-model baseline regenerated 2026-09-02, 48 rows each, first failing test in metadata), code-bench-v4__opus-5.5-bedrock (added 2026-09-22 via Amazon Bedrock), code-bench-v4__sonnet-5.5-bedrock (added 2026-09-28 via Amazon Bedrock) and code-bench-v4-july__* (July 25 outputs re-scored); repository code-benchmarks (harness, tests, outputs, this report). Prepared by Andrés Espitia with Claude Code.